

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ models evaluated across Agentic · Frontier · Safety · Multimodal

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

Coding & Tool Use

- SWE Atlas
- SWE-Bench Pro
- MCP Atlas
- VisualToolBench

Frontier

Reasoning & Knowledge

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- SciPredict

Multimodal

Audio & Vision


- AudioMultiChallenge
- VISTA
- VisualToolBench

Safety

Alignment & Honesty

- MASK
- Fortress
- PropensityBench
- Professional Reasoning


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro



Benchmark	#1 Model	Score	±
SWE Atlas – Test Writing	GPT-5.4-xHigh	44.36	6.04
SWE Atlas – Codebase QnA	GPT-5.4-xHigh	40.80	5.10
SWE-Bench Pro (Public)	GPT-5.4-pro	59.10	3.56
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	6.07

MCP Atlas #1: Muse Spark (78.3)


GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%



Insight: EnigmaEval (24%) and SciPredict (25%) show frontier models remain far from solving complex multi-step reasoning.

Gemini Models Sweep Audio and Vision-Language Benchmarks

Multimodal Benchmarks: Top Scores



AudioMultiChallenge (Text Output)

Gemini-3-pro-preview leads at 54.65 ± 4.57
Gemini-2.5-pro (46.90) and Gemini-2.5-flash (40.04) complete a clean sweep of the top 3.

AudioMultiChallenge (Audio Output)

Gemini-3.1-flash-live-preview leads at 36.06 ± 4.41
GPT-realtime-1.5 follows at 34.73.



VISTA (Vision-Language)

Gemini 2.5 Pro Experimental leads at 54.65 ± 1.46
GPT-5.4-pro close behind at 53.89 ± 2.02.

VisualToolBench

GPT-5.4 edges Gemini-3.1-pro
29.17 vs 28.97 — tight race in visual tool use.

Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low



Contrast: MASK honesty scores approach ceiling (96%+) while Fortress national-security scores remain in the single digits.
PropensityBench (lower=better): o3 leads at 10.50 ±0.60 | RLI: Claude Opus 4.6 (CoWo) at 4.17

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Domain	Benchmark	OpenAI	Google	Anthropic	Muse Spark	Other
Agentic	SWE Atlas – Test Writing	★				
Agentic	SWE Atlas – Codebase QnA	★				
Agentic	MCP Atlas				★	
Agentic	SWE-Bench Pro (Public)	★				
Agentic	SWE-Bench Pro (Private)			★		
Agentic	VisualToolBench	★				
Frontier	Humanity's Last Exam	★				
Frontier	HLE (Text Only)	★				
Frontier	EnigmaEval	★				
Frontier	MultiChallenge				★	
Frontier	MultiNRC	★				
Frontier	SciPredict		★			
Multimodal	AudioMultiChallenge (Text)		★			
Multimodal	AudioMultiChallenge (Audio)		★			
Multimodal	VISTA		★			
★ = #1 ranked model on that benchmark Muse Spark appears as a new entrant with #1 on MCP Atlas and MultiChallenge						
Safety	SAFE (Honesty)			★		

Key Takeaways: No Single Model Wins Everything — ~~Specialization~~ Is the New Frontier

1. GPT-5.4 leads agentic coding and frontier reasoning, but does not dominate multimodal or safety benchmarks.
2. Gemini sweeps audio and vision-language leaderboards, holding all top-3 positions on AudioMultiChallenge (Text Output).
3. Claude leads safety and honesty with 96%+ MASK scores and #1 on both professional reasoning benchmarks (Legal & Finance).
4. Benchmarks remain far from saturation: HLE (44%), EnigmaEval (24%), and SciPredict (25%) show substantial headroom.
5. Muse Spark emerges as a strong new competitor, ranking top 3 on 6+ benchmarks — including #1 on MultiChallenge at 75.5%.

Conclusion: The frontier is fragmented. Labs must choose where to specialize — or invest broadly to close capability gaps.